

Stratis Software Design: Version 2.0.0*

Andy Grover <agrover@redhat.com>

Last modified: 11/08/2019

Contents

I	Background	3
1	Problem Statement	3
1.1	Goal: Bring advanced features to users in a simpler form	4
1.2	Proposal: Implement a hybrid Volume Managing Filesystem	4
1.3	Requirements	4
II	Solution Overview	6
2	Introduction	6
3	Stratis and the Linux storage stack	6
4	Conceptual Model	7
4.1	Blockdevs, pools, and filesystems	7
4.2	Attributes and features of a pool	7
5	Scalability and Performance Considerations	7
III	Implementation	7
6	Software Components	7
7	User Experience	8
7.1	Known shortcomings	8
8	D-Bus Programmatic API	9
8.1	Overview	9
8.2	D-Bus Access Control	10
8.2.1	Security Policy	10
8.2.2	Prevent Spoofing	10

*This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International License.

9	Internals	10
9.1	Data Tier Requirements	10
9.2	Data Tier	10
9.2.1	Layer 0: Blockdevs	10
9.2.2	Layer 1: Integrity (optional) ¹	11
9.2.3	Layer 2: Redundancy (optional) ²	11
9.2.4	Cache Tier (optional)	12
9.2.5	Layer 3: Flex	12
9.2.6	Layer 4: Thin Provisioning	12
9.2.7	Layer 5: Thin Volumes	13
9.2.8	Layer 6: Encryption (optional) ³	13
9.2.9	Layer 7: Filesystem	13
9.3	Data Tier Metadata	13
9.3.1	Requirements	13
9.3.2	Conventions	14
9.3.3	Design Overview	14
9.3.4	BlockDev Data Area (BDA)	15
9.3.5	Metadata Area (MDA)	17
9.3.6	Metadata Volume (MDV)	19
9.3.7	The MDA and Very Large Pools	19
9.3.8	Metadata and Recovery	19
9.4	Cache Tier	19
9.4.1	Requirements	20
9.5	Cache Tier Metadata	20
9.5.1	Cache Tier Metadata Requirements	20
10	Implementation Details	20
10.1	'stratis' command-line tool	20
10.2	stratisd	20
10.3	device-mapper names	21
10.3.1	Naming convention Requirements	21
10.3.2	Naming Convention	21
10.3.3	Per-Layer Naming Definitions	21
10.4	device-mapper minimum version	22
10.5	OS Integration: Boot and initrd	22
10.6	OS Integration: udev	22
10.7	OS Integration: /dev entries	23
10.8	Partial Pool Activation	23
10.9	Snapshots	23
10.10	Backstore Internals	23
10.10.1	Demand-based allocations	25
10.10.2	Common Layer API	25
10.11	Operation States	25
10.11.1	Pool States	25
10.11.2	Pool Extend States	25
10.11.3	Pool Space States	26
10.12	Licensing	26
IV	Development Plan	26

11 Delivery of Features	26
11.1 Stratis version 0.1	26
11.2 Stratis version 0.5	27
11.3 Stratis version 1.0	27
11.4 Stratis version 2.0	28
11.5 Stratis version 3.0	28
11.6 Stratis version 4.0	28
A Initial options for adapting existing solutions	29
A.1 Extending an existing project	29
A.1.1 SSM	29
A.1.2 LVM2	29
A.2 Building upon existing projects	30
A.2.1 XFS	30
A.2.2 device-mapper	30
A.2.3 LVM2	30
A.3 Conclusions	32

Asking Questions and Making Changes to this Document

This document can be found in the stratis-docs repo, and is written using LyX 2.3.0. Please use the mailing list (stratis-devel@lists.fedorahosted.org), or open an issue on GitHub, for any questions or issues that arise.

Executive Summary

Stratis is a new tool that meets the needs of Red Hat Enterprise Linux (RHEL) users (among others) calling for an easily configured, tightly integrated solution for storage that works within the existing Linux storage management stack. To achieve this, Stratis prioritizes a straightforward command-line experience, a rich API, and a fully automated, externally-opaque approach to storage management. It builds upon elements of the existing storage stack as much as possible, to enable delivery within 1-2 years. Specifically, Stratis initially plans to use device-mapper and the XFS filesystem, but may incorporate other technology in the future.

Part I

Background

1 Problem Statement

Linux has gained many storage-related features over the years, but each of these features has required the user to manage the configuration of these features in a layered, additive manner. Genuinely new and useful features such as thin provisioning, RAID, and multipath are dependent on the user correctly configuring many different layers via different tools to achieve a complete result. Furthermore, since each layer's configuration tool only has a command-line interface (CLI), higher-level management tools must each construct input and parse the human-oriented output for each these layers' CLI. This causes a waste of effort and opportunity for bugs, as each higher-level tool builds its own internal API for the feature on top of the lower level tool's CLI.

1.1 Goal: Bring advanced features to users in a simpler form

Linux storage features are modular and stackable. This promotes flexibility and allows independent development efforts, but leads to a huge number of possible configurations. This requires the user manage the stack because there's not enough commonality to enable effective automation.

But really, there *is* a single configuration that can work for most use cases. By assuming a fixed layering of storage features (some perhaps optional), we enable software to effectively manage these on behalf of the user.

Automated management then leads to less administrative burden placed on the user. The user still specifies resources, desired features, and results – what hardware resources to use, what features to enable, how storage should be logically presented – using a smaller number of concepts with well-defined relations. Software manages the rest, and handles most runtime issues without user involvement.

1.2 Proposal: Implement a hybrid Volume Managing Filesystem

In the past ten years, *volume-managing filesystems* (VMFs) such as ZFS and Btrfs have come into vogue and gained users, after being previously available only on other UNIX-based operating systems. These incorporate what would be handled by multiple tools under traditional Linux into a single tool. Redundancy, thin provisioning, volume management, and filesystems become features within a single comprehensive, consistent configuration system. Where a traditional Linux storage stack exposes the layers of block devices to the user to manage, VMFs hide everything in a *pool*. The user puts raw storage in the pool, the VMF manages the storage in the pool, providing the features the user wants, and allows the user to create filesystems from the pool without being concerned with the details.

Unfortunately, existing VMFs aren't easily used on enterprise Linux distributions like RHEL. ZFS isn't an option RHEL can embrace due to licensing, Ubuntu notwithstanding. Btrfs has no licensing issues, but maintaining up-to-date support for it in enterprise kernels proved difficult.

We can see from the many developer-years of effort that have gone into these two projects that writing a VMF is a tremendous, time-consuming undertaking. We also can hear our users demanding their features and ease of use.

Rather than writing a new VMF from scratch, Stratis proposes to satisfy VMF-like requirements by managing existing technologies on behalf of the user, so that users can manage their storage using high-level concepts like “pool” and “filesystem”, and remain unconcerned with the more complex details under the covers.

This is also a chance to learn from the benefits and shortcomings of existing solutions. We should not just copy ZFS. ZFS is now fifteen years old and the storage landscape has changed since its design. We seek to satisfy the same needs that ZFS does, but also integrate more tightly into today's increasingly automated storage management solutions that span the data center as well as the local machine. This is made possible by a hybrid, userspace-based approach.

1.3 Requirements

1. *Make features easier to use in combination with each other*: thin provisioning, snapshots, integrity, redundancy, multipath, encryption, hardware reconfiguration, monitoring, and a caching tier
2. Simple and comprehensive command-line interface
 - (a) Simple
 - i. Single way to do things
 - ii. Do not expose internal implementation details. Gives Stratis more implementation freedom, and of little value since internals are too complex to make manual user repairs practical

- iii. User typically will not use on a daily basis
 - A. Consistent commands that a user can guess at, and probably be right
 - B. Require explicitness from the user for potentially data-losing operations, such as giving a “-force” option.
 - (b) Comprehensive
 - i. User must master only one tool
 - ii. Helps user learn: if task not possible through tool, it must not be worth doing (or a good idea)
- 3. Programmatic language-neutral API for higher-level management tool integration
 - (a) A clear next step for users after hitting the limitations of scripting the CLI
 - (b) Encourages tight integration and use of all features by higher-level tools
- 4. Event-driven monitoring and alerts
 - (a) Monitoring and alert messages expressed in terms of Stratis user-visible simple concepts, not implementation details
 - (b) Low CPU/memory overhead to monitoring
 - (c) Only alert when action really is needed
 - (d) Fail gracefully if alerts are unheeded
- 5. Eliminate manual resizing of filesystems
 - (a) Numerous problem reports throughout the years indicate that resizing filesystems is an area where users feel unease, due to potential data loss if a mistake is made. No real reason to require the user do this any more.
 - (b) Simpler for DevOps
 - (c) Makes storage “demand-allocated”, similar to virtual memory. Current technology allows us to manage a filesystem’s actual usage up (growfs) or down (thin provisioning).
- 6. Initrd-capable
 - (a) Allows root fs, all other filesystems except /boot to use Stratis. Needed for ease of use
 - (b) Limited environment – no Python or Dbus – but can use device-mapper
- 7. Bootable (planned – see 11.6)
 - (a) Feels like a “real” filesystem if no secondary filesystem is needed for boot
 - (b) Enables Stratis features to be used by system image, e.g. booting from a snapshot, and allowing /boot to grow
 - (c) Requires explicit support in bootloader (Grub2)
 - (d) device-mapper not available
- 8. Adaptable to emerging storage technologies
 - (a) Persistent memory
 - i. Block-appearing pmem can be used by Stratis
 - ii. byte-addressable pmem see 11.6
- 9. Implementable in 1-2 years

(a) We're already behind, waiting another 10 years isn't an option

Part II

Solution Overview

2 Introduction

Stratis is a local storage solution that lets multiple logical filesystems share a pool of storage that is allocated from one or more block devices. Instead of an entirely in-kernel approach like ZFS or Btrfs, Stratis uses a hybrid user/kernel approach that builds upon existing block capabilities like device-mapper, existing filesystem capabilities like XFS, and a user space daemon for monitoring and control.

The goal of Stratis is to provide the conceptual simplicity of volume-managing filesystems, and surpass them in areas such as monitoring and notification, automatic reconfiguration, and integration with higher-level storage management frameworks.

3 Stratis and the Linux storage stack

Stratis simplifies many aspects of local storage provisioning and configuration. This, along with its API, would let projects dependent on configuring local storage do so much more easily.

For example, installing the OS to a Stratis pool using Anaconda. After selecting the disks to use for the pool, the first benefit would be the complex flow around sizing of filesystems could be omitted. Second, since Stratis has an API, Anaconda could use it directly, instead of needing work in Blivet to build an API on top of command line tools.

Other management tools like Cockpit, virtualization products like RHEV, or container products like Atomic would find it much simpler and less error-prone to use storage and snapshots with Stratis, for the same two reasons: don't need to worry about per-filesystem sizing (only that the pool has enough "backing store"); and the API, which allows better tool-to-tool integration than using CLI programmatically.

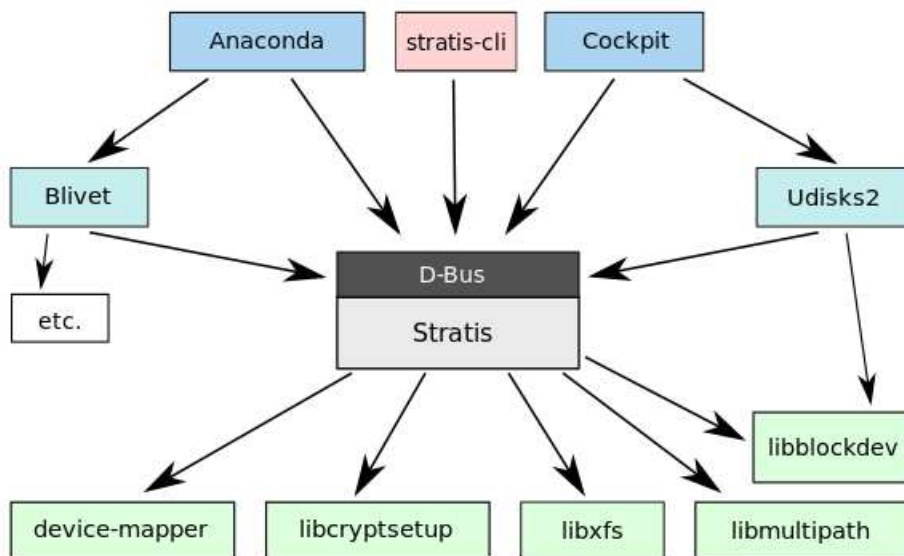


Figure 1: Future Stratis Position in the Linux Storage Management Stack

There are existing libraries that handle CLI-to-API for both Anaconda and Cockpit. These could be extended to support Stratis, or not.

4 Conceptual Model

4.1 Blockdevs, pools, and filesystems

Stratis's conceptual model consists of *blockdevs*, *pools*, and *filesystems*. A pool is created from one or more blockdevs (block devices), and then filesystems are created from the pool. Filesystems are mountable hierarchical collections of files that allocate backing storage from the pool as it is needed. The key difference between a Stratis filesystem and a conventional Unix filesystem is that Stratis filesystem sizing and maintenance are not managed by the user, but by Stratis.

4.2 Attributes and features of a pool

A pool is created with an initial set of one or more blockdevs. Blockdevs may also be added after the pool is created. Blockdevs may eventually be removable from a pool, if certain preparations are performed and conditions are met. The pool's primary collection of blockdevs is called the *data tier*.

A pool also optionally has a *cache tier* that uses a separate collection of faster blockdevs to improve performance instead of increase the pool's capacity.

Both tiers have their own independent data redundancy level. which may be none, raid1, raid5, raid6, or raid10 (the default is none).

Since a single system may have multiple pools, each pool has a name, as does each filesystem within a pool. These are both settable by the user. Blockdevs, pools, and filesystems also have UUIDs, which are not settable by the user.

Stratis supports large numbers of blockdevs and up to 2^{24} filesystems per pool. However, practical limits on these values may compel users to restrict themselves to smaller numbers of blockdevs and filesystems.

A new filesystem is either a new empty filesystem or a snapshot of an existing filesystem within the pool. Stratis currently does not distinguish between snapshots and filesystems. (May change.)

5 Scalability and Performance Considerations

Stratis doesn't optimize performance within its data tier, instead focusing there on flexibility and integrity. Improved performance is the job of caching tier, or perhaps building the pool using blockdevs with higher IOPs, such as SSDs.

Part III

Implementation

6 Software Components

Stratis consists of a command-line tool, *stratis*, and a service, *stratisd*.

stratis implements the command-line interface, and converts commands into D-Bus API calls to *stratisd*.

stratisd implements the D-Bus interface, and manages and monitors Stratis internal pool blockdevs, as described below. It is started by the system and continues to run as long as Stratis pools or blockdevs are present in the system.




Figure 2: Stratis Architecture

stratisd includes a simulator engine. The simulator engine is purely computational and does not affect the environment, although it does communicate over the D-Bus.

Figure 2 shows the basic Stratis architecture.

7 User Experience

Stratis has a command-line tool that enables the administrator to create a Stratis pool from one or more blockdevs, and then allocate filesystems from the pool.

See reference implementation at <https://github.com/stratis-storage/stratis-cli> for the most up-to-date status of the CLI design.

This component is not required to be installed, in cases such as an appliance where a higher-level application such as Cockpit or Ansible uses the D-Bus API directly.

7.1 Known shortcomings

Stratis' goal is to hide the complexity of its implementation from the user, but by using a reuse/layering approach to its implementation, there will be places where Stratis' implementation

details will peek through. This could cause user confusion, and also could threaten Stratis integrity if the user makes changes.

- For Stratis filesystems, 'df' will report the current used and free sizes as seen and reported by XFS. This is not useful information, because the filesystem's actual storage usage will be less due to thin provisioning, and also because Stratis will automatically grow the filesystem if it nears XFS's currently sized capacity.
- Users should not try to reformat or reconfigure XFS filesystems that are managed by Stratis. Stratis has no way to enforce this or warn the user to avoid this, other than in the documentation.
- Stratis will use many device-mapper devices, which will show up in 'dmsetup' listings and /proc/partitions. Similarly, 'lsblk' output on a Stratis system will reflect Stratis' internal workings and layers.
- Stratis requires a userspace daemon, which must remain running at all times for proper monitoring and pool maintenance.

8 D-Bus Programmatic API

The Stratis service process exposes a D-Bus interface, for other programs to integrate support for Stratis. This is considered the primary Stratis interface. The command-line tool uses the D-Bus API.

8.1 Overview

The D-Bus API is part of `stratisd`. It is a thin layer that receives messages on the D-Bus, processes them, transmits them to the Stratis engine, receives the results from the engine, and returns the result to the invoker of the API. When processing method calls, its responsibilities are confined to:

- Receiving arguments and verifying that they conform to the signature of the invoked method.
- Transforming method arguments received on the D-Bus to arguments of the appropriate type to be passed to engine methods.
- Converting tuple arguments used to represent non-mandatory arguments to values which inhabit the Rust Option type.
- Invoking the appropriate engine methods and capturing their return values.
- Marshalling the appropriate return values to place on the D-Bus along with the return code and message.
- Adding or removing objects from the D-Bus tree.

The D-Bus API is implemented using the `dbus-rs` library⁴.

The Stratisd D-Bus API Reference Manual contains a description of the API.

⁴<https://github.com/diwic/dbus-rs>

8.2 D-Bus Access Control

8.2.1 Security Policy

Most stratisd D-Bus methods require root permissions. However, listing operations do not; these can be done by an unprivileged user. The default permissions are specified in the policy file, `stratisd.conf`, included in the stratisd distribution. These defaults permit all actions by root users but restrict unprivileged users to read-only actions. Systems administrators can adjust permissions by editing the stratisd D-Bus policy files.

8.2.2 Prevent Spoofing

9 Internals

Stratis internals aim to be opaque to the user. This allows its implementation maximum flexibility to do whatever it needs in Stratis version 1, as well as to be extended in later versions without violating any user-visible expectations.

9.1 Data Tier Requirements

The data tier of Stratis must manage blockdevs on behalf of the user to provide the following:

1. Managed filesystems that consume only as much space as the files they contain
2. Fast snapshots of existing filesystems
3. The ability to add (and eventually remove) individual blockdevs to grow the available space available to filesystems
4. User-selectable redundancy level (per-pool granularity)⁵
5. Integrity checking⁶
6. Encryption⁷

9.2 Data Tier

The data tier achieves these requirements by layering device-mapper (DM) devices on top of the pool's blockdevs. The topmost layer consists of thin devices allocated from a thinpool. Stratis initializes these thin devices with a filesystem, and manages the DM devices and filesystems to meet the above requirements.

9.2.1 Layer 0: Blockdevs

This layer is responsible for discovering existing blockdevs in a pool, initializing and labeling new blockdevs unambiguously as part of the pool, setting up any disk-specific parameters, and storing pool metadata on each blockdev. The minimum blockdev size Stratis will accept is 1 GiB. Blockdevs may be encrypted or unencrypted.

⁵in a future version

⁶in a future version

⁷in a future version

Encryption If the blockdevs are encrypted, the following conditions will hold:

- Each blockdev will be encrypted with a distinct and randomly generated MEK (Media Encryption Key).
- All blockdevs in a pool will be encrypted, or all blockdevs in a pool will be unencrypted. Stratis will support mixed usage for pools, where some pools are encrypted and others are not.
- A distinct KEK (Key Encryption Key) for each pool will be supported, although it will be possible to use the same KEK for every pool, if desired.
- The user will be required to specify the keyfile containing the KEK at pool creation time via the CLI. Encryption of an already existing pool will not be supported.
- Re-encryption will not be supported initially; this functionality will be considered and added in a later step.
- If a pool was encrypted on creation, then all blockdevs added to the data tier will be automatically encrypted with a randomly generated MEK and the pool's designated KEK.
- On any encrypted blockdev, the Stratis metadata will itself be encrypted; it will be inaccessible until the encrypted blockdev is opened.
- For encrypted blockdevs, the minimum blockdev size will apply to the opened device. Thus the entire size of the encrypted blockdev will be required to be more than 1 GiB.
- The use of a cache tier and of encryption will be mutually exclusive. If a pool is encrypted, an attempt to add a blockdev to the cache tier will be rejected.

The above conditions require an implementation which makes use of the LUKS2 header. In particular, for a blockdev to be decrypted with its pool-specific KEK, it will be necessary to include information within the device header which allows the device to be identified as belonging to a particular Stratis pool. The LUKS2 header includes support for LUKS tokens, which will allow `stratisd` to identify the pool to which a device belongs; the LUKS1 header does not.

9.2.2 Layer 1: Integrity (optional)⁸

This layer uses the `dm-integrity` target to enable the detection of incorrect data as it is read, by using extra space to record the results of checksum/hash functions on the data blocks, and then comparing the results with what the blockdev actually returned. This will enable Stratis to detect data corruption when the pool is non-redundant, and to repair the corruption when the pool is redundant. It should also be possible to use DIF⁹ information if present.

9.2.3 Layer 2: Redundancy (optional)¹⁰

A Stratis pool may optionally be configured to spread data across multiple physical disks, so that the loss of any one disk does not cause data loss. Stratis uses conventional RAID technology (1, 5, 6, 10, 1E¹¹) as specified, and converts Layer 0 blockdevs into a smaller-sized amount of storage with the specified raid properties.

Since Stratis supports more blockdevs than are RAID-able (generally 8 or fewer is best for a raidset¹²), and differently-sized blockdevs, a redundant Stratis pool may contain multiple raid sets (all of the same global type). Depending on layout, there may be some amount of space

⁸in a future version

⁹See <https://oss.oracle.com/projects/data-integrity/>

¹⁰in a future version

¹¹See 'linux/Documentation/device-mapper/dm-raid.txt' for more info

¹²More are supported by DM, but too many increase the likelihood of individual failures.

in a pool's blockdevs that cannot be used because it cannot be used in a RAID set. Stratis will intensively manage raidsets, extending them across newly added blockdevs or creating new raidsets; and handling the removal of blockdevs¹³. Stratis may use dm-raid's reshape capabilities when possible, although this changes the stripe size and could cause issues.

Stratis cannot support redundancy with a single disk, but we may wish to reserve the small space for raid metadata and other uses even on one-disk Stratis pools. This will allow the pool to be made redundant (in the version when we support this) without encountering ugly edge cases.

9.2.4 Cache Tier (optional)

If present, the Cache Tier sits directly underneath the Flex Layer. Its structure is similar to the lower levels of the Data Tier. See 9.4 for more.

9.2.5 Layer 3: Flex

Whether blockdevs are part of raidsets or used directly, pools need to cope with the addition or removal¹⁴ of them.

Stratis allows adding a blockdev to an existing pool, and using it to grow the pool's allocated space.

Removal of a blockdev involves calculating if a blockdev can be removed from the pool with no effect, removed with loss of redundancy, not removed without data loss, or if removal would be possible after a data movement and reorganization step. The data-movement capability could also be used in the case of a failing disk.

Marking a newly-added blockdev as 'spare' would keep it in reserve so a failing disk can be immediately removed and replaced without possibly exceeding total pool capacity in the interim.

The flexibility layer contains four linear DM devices made up of segments from lower-level devices. The first two devices will be used by Layer 4 (Thin Provisioning) as metadata and data devices. The flex layer will track what lower-level devices these are allocated from, and allow the two devices to grow as needed.

The third linear DM device is a spare metadata device to be used in the case that the metadata device requires offline repair. It will not usually be instantiated on the system, but guarantees there is room if needed. This device's size tracks the size of the metadata device, both as initially allocated, and as the metadata device is extended.

The fourth and final linear DM device is used for the Metadata Volume (MDV, see 9.3.6). The MDV is used to store metadata about upper layers, layer five and above.

All devices in this layer may be built on L0, L1, or L2 devices, depending on configuration.

The initial sizes of all flex layer devices should be chosen to allow an entire pool to fit within a single blockdev of the minimum size (1 GiB).

9.2.6 Layer 4: Thin Provisioning

The two linear targets from L3 are used as metadata device and data device for a DM thinpool device. The thinpool device implements a copy-on-write (CoW) algorithm, so that blocks in the data device are only allocated as needed to back the thin volumes created from the thinpool.

Stratis manages the thinpool device by extending the two L3 subdevices when either runs low on available blocks. If the pool approaches a point where the pool no longer has empty lower-level space to extend onto, Stratis alerts the user and takes action to avoid data corruption. Actions could include switching filesystems to read-only mode, running `fstrim`, or progressively throttling writes.

¹³in a future version

¹⁴in a future version

9.2.7 Layer 5: Thin Volumes

Stratis creates thin volumes from the thin pool. It will automatically give a new volume a default size, format it with a filesystem, and make it available to the user.

Stratis also enables creating a new volume as a read/write snapshot of an existing volume. Although the underlying implementation does not require maintaining the relation between a snapshot and its origin, Stratis records this relation in its metadata. This relation may be of use to users who may, for example, use snapshots for backups and may make use of the origin information to identify a particular backup snapshot to restore from.

9.2.8 Layer 6: Encryption (optional)¹⁵

Stratis may eventually enable per-filesystem encryption between the thin device and the filesystem.

9.2.9 Layer 7: Filesystem

Stratis monitors each filesystem's usage against its capacity and automatically extends them online without user intervention. Extending involves changing the thin dev's logical size, and then using a tool such as `xfs_growfs` to grow the filesystem. Stratis also periodically will run `fstrim` to reclaim unused space during idle periods. Idle periods will be found based upon current and historical system I/O activity levels.

9.3 Data Tier Metadata

Stratis must track the blockdevs that make up the data tier of the pool (L0), integrity parameters (L1), the raidsets that are created from the data blockdevs (L2), the three linear targets that span the L2 devices (L3), the thinpool device (L4) and the attributes of the thin devices (L5) and filesystems created from the thinpool (L7).

9.3.1 Requirements

1. Uniquely identify a blockdev as used by Stratis, which pool it is a member of, and parameters needed to recreate all layers
2. Detect incomplete or corrupted metadata and recover via second copy
3. Allow for blockdevs being expanded underneath Stratis
4. Redundant on each blockdev to tolerate unreadable sectors¹⁶
5. Redundant across blockdevs to handle missing or damaged members. Can provide metadata of missing blockdevs
6. Handle thousand+ blockdevs in a pool
7. Handle million+ filesystems in a pool and updates without writing to each blockdev
8. Extensible/upgradable metadata format

¹⁵in a future version

¹⁶Recovery from accidental start-of-blockdev overwrite by placing a second copy at the end of the disk was also considered, but raised other issues that outweighed its benefit.

9.3.2 Conventions

Sectors are 512 bytes in length¹⁷.

All UUIDs are written as un-hyphenated ASCII encodings of their lower-case hexadecimal representation¹⁸.

9.3.3 Design Overview

Stratis metadata is in three places:

1. Blockdev Data Area (BDA)
 - (a) Signature Block within Static Header
 - (b) Metadata Area (MDA)
2. Metadata Volume (MDV)

(Specific DM targets such as raid, integrity, and thinpool also place their own metadata on disk.)

Information on levels 0-4 is duplicated across all blockdevs within a on-disk metadata format called the Blockdev Data Area (BDA). The BDA consists of a binary Signature Block, and the Metadata Area (MDA), which stores information in a text-based JSON format. Both the binary and text-based portions of the BDA define redundancy and integrity-checking measures.

The Metadata Volume (MDV) stores metadata on Layers 5 and up in a conventional block device and filesystem that is part of the Flex layer. Choosing to split overall metadata storage into two schemes allows upper layers' metadata to be free of limitations that would apply if a single scheme was used. For example, on-disk metadata formats find it hard to support runtime size extension, may keep redundant copies to ensure reliability, and aggressively check for corruption. This can work well with small amounts of data that is infrequently changed, but has trouble as data grows, or we wish to do updates in-place.

Upper-level metadata can achieve redundancy and integrity by building on the pre-existing lower layers, and work under looser restrictions around updating in place, and the total size to which it may grow. It can reuse an existing, well-tested solution for solving data organization and storage issues – a general-purpose filesystem.

¹⁷Historically this is the minimum storage unit of a hard drive. Many Linux kernel APIs assume this value is constant (as does this document), and use another term such as 'block size' for dealing with cases where the minimum storage unit is different.

¹⁸UUIDs are 128-bit values and therefore require only 16 bytes to represent their numeric value. However, since each ASCII value requires a byte, and the hexadecimal representation of an 128-bit value requires 32 hexadecimal digits, the chosen encoding requires 32 bytes.

9.3.4 BlockDev Data Area (BDA)




Figure 3: BDA format

The BDA consists of a fixed-length Static Header of sixteen sectors, which contains two copies of the Signature Block; and the metadata area (MDA), whose length is specified in the Signature Block. These are written to the beginning of the blockdev as described below.

Stratis reserves the first 16 sectors of each blockdev for the Static Header. When initializing or modifying the Signature Block, identical data is written to locations 1 and 2.

Static Header

sector offset	length (sectors)	contents
0	1	unused
1	1	Signature Block location 1
2	7	unused
9	1	Signature Block location 2
10	6	unused

Signature Block

byte offset	length (bytes)	description
0	4	IEEE CRC32C of signature block (bytes at offset 4 length 508)
4	16	Stratis signature: '!Stra0tis\x86\xff\x02^\x41rh'
20	8	Device size in 512-byte sectors (u64)
28	1	Signature Block version (u8) (value = 1)
29	3	unused
32	32	UUID of the Stratis pool
64	32	UUID of the blockdev
96	8	sector length of blockdev metadata area (MDA) (u64)
104	8	sector length of reserved space (u64)
112	8	flags (u64)
120	8	initialization time: UNIX timestamp (seconds since Jan 1 1970) using UTC (u64)
128	384	unused

- No flags are yet defined, so 'flags' field is zeroed.
- All 'unused' fields are zeroed, and are reserved for future use.
- If not zero, blockdev metadata area length (offset 96) must be a number divisible by four of at least 2032.
- The BDA is followed immediately by *reserved space*, whose size is specified in the signature block (offset 104).
- Minimum length of BDA (static header and MDA) plus Reserved Space is 2048 sectors (1 MiB).
- When a blockdev is removed from a pool, or is part of a pool that is destroyed, Stratis wipes the Static Header.
- The purpose of the unused sectors is twofold. First, placing the Signature Block copy locations in two separate 4K blocks helps to prevent a single bad write operation on 4K-block disks from corrupting both copies. Second, using a single sector for the Signature Block helps to minimize the likelihood of corruption on disks with 512 byte blocks.
- Each time that Stratis writes one or both Signature Block locations, it also zeroes the unused sectors that share the same 4K block.

The MDA is divided into four equal-size regions, numbered 0-3. When updating metadata, identical data is written to either the odd (1 and 3) or even (0 and 2) regions, chosen by examining the timestamps and overwriting the older of two pairs.

Each MDA region's update consists of a fixed-length MDA Region Header, followed by variable-length JSON data.

MDA Region Header

byte offset	length (bytes)	description
0	4	IEEE CRC32 covering remainder of MDA header
4	4	IEEE CRC32 covering JSON data
8	8	length of JSON data in bytes (u64)
16	8	UNIX timestamp (seconds since Jan 1 1970) using UTC (u64)
24	4	nanoseconds (u32)
28	1	Region Header version (u8) (value = 1)
29	1	Variable-length metadata version (u8) (value = 1)
30	2	unused
32	variable	JSON data

- Metadata updates write to the older of the odd or even MDA regions. This is determined by lowest timestamp, and then lowest nanoseconds if timestamps are equal.
- MDA updates include the MDA Header, which includes the current time. However, if using the current time would not result in the update having the latest time across all MDA regions on all blockdevs in the pool, instead use a time of one nanosecond later than the latest MDA region time across all blockdevs.
- The procedure for updating metadata is:
 1. Determine which regions in the MDA to use (odd or even) as described above.
 2. Write MDA header and JSON data to the first MDA region (0 or 1)
 3. Perform a Flush/FUA
 4. Write MDA header and JSON data to the second MDA region (2 or 3)
 5. Perform a Flush/FUA
 6. Repeat for additional blockdevs. Also see 9.3.7
- Multiple blockdevs being updated with the same metadata must write identical data to each MDA region, but which regions (odd or even) is used may vary, if the blockdevs have received differing numbers of metadata updates over time.

9.3.5 Metadata Area (MDA)

The MDA contains a JSON object that represents the pool's overall configuration from L0 to L4.

Top level objects:

key	JSON type	required	description
name	string	y	the name of the pool
backstore	object	y	the block devices in the pool
flex_devs	object	y	layout of the data and metadata linear devices
thinpool_dev	object	y	parameters of the thinpool device

backstore: An object describing the data tier and the cache tier.

key	JSON type	required	description
data_tier	object	y	the block devices in the data tier
cap	object	y	the cap device, from which segments are allocated to the flex layer
cache_tier	object	n	the block devices in the cache tier

data_tier: An object describing the data tier.

key	JSON type	required	description
blockdev	object	y	Settings and mappings describing block devices that make up the tier
integrity	object	n	TBD settings and mappings associated with integrity support
raid	object	n	TBD settings and mappings associated with redundancy (RAID) support
vdo	object	n	TBD settings and mappings associated with VDO

blockdev: An object describing physical block devices that make up the tier.

key	JSON type	required	description
devs	array	y	an array of base_block_dev objects
allocs	array	y	an array of arrays of base_dev objects

base_dev: An object describing an allocation from a block device.

key	JSON type	required	description
parent	string	y	UUID of the device the segment is created from
start	integer	y	the starting sector offset within the parent device
length	integer	y	the length in sectors of the segment

base_block_dev: An object describing a block device in the lowest layer.

key	JSON type	required	description
uuid	string	y	The UUID of the block device, as recorded in its Signature Block
user_info	string	n	user-provided information for tracking the device
hardware_info	string	n	uniquely identifying information for the blockdev, such as SCSI VPD83 or serial number

integrity: (TBD) Settings and mappings associated with integrity support.

raid: (TBD) Settings and mappings associated with RAID support.

vdo: (TBD) Settings and mappings associated with VDO support.

cap: An object describing a view of the top-level linear device provided by the backstore to the flex layer.

key	JSON type	required	description
allocs	array	y	an array of pairs of integers representing the starting offset and length of an allocation in sectors

cache_tier: An object describing the cache tier. Identical format to **data_tier** except VDO layer is not supported.

flex_devs: An object with four keys that define the linear segments that make up each device in the Flex layer.

key	JSON type	required	description
meta_dev	array	y	an array of pairs of integers representing the starting offset and length of an allocation in sectors that make up the metadata volume (MDV)
thin_meta_dev	array	y	an array of pairs of integers representing the starting offset and length of an allocation in sectors that make up the thin metadata device
thin_meta_dev_spare	array	y	an array of pairs of integers representing the starting offset and length of an allocation in sectors that make up the thin metadata spare device
thin_data_dev	array	y	an array of pairs of integers representing the starting offset and length of an allocation in sectors that make up the thin data device

thinpool_dev: An object that defines properties of the thinpool device.

key	JSON type	required	description
data_block_size	integer	y	the size in sectors of the thinpool data block

9.3.6 Metadata Volume (MDV)

The Metadata Volume is formatted with an XFS filesystem that is used by Stratis to store information on user-created thin filesystems (L5-L7). This information is stored in the filesystem in a TBD format, maybe either an individual file-based scheme, or SQLite database.

9.3.7 The MDA and Very Large Pools

Stratis pools with very large numbers of blockdevs will encounter two issues. First, updating the metadata on all blockdevs in the pool may become a performance bottleneck. Second, the default MDA size may become inadequate to contain the information required.

To solve the first issue, Stratis caps the number of blockdevs that receive updated metadata information. A reasonable value for this cap might be in the range of 6 to 10, and should try to spread metadata updates across path-independent blockdevs, if this can be discerned, or randomly. This limits excessive I/O when blockdevs are added or removed from the pool, but maximizes the likelihood that up-to-date pool metadata is retrievable in case of failure.

To solve the second issue, Stratis monitors how large its most recent serialized metadata updates are, and increases the size of MDA areas on newly added devices when a fairly low threshold – %50 – is reached in comparison to the available MDA region size. This ensures that by the time sufficient blockdevs have been added to the pool to be in danger of serialized JSON data being too large, there are enough blockdevs with enlarged MDA space that they can be used for MDA writes.

9.3.8 Metadata and Recovery

Bad things happen.

In order to recover from disk errors, Stratis uses CRCs over the critical L0-L4 metadata, and writes duplicate copies to a single blockdev, as well as across multiple blockdevs, when possible. It takes this approach – copies – rather than a mechanism that might make it possible to partially repair corrupted metadata for three reasons:

1. This metadata is relatively small – it concerns disks and raidsets, of which the pool will have only a small number, so having multiple entire copies isn't terribly wasteful.
2. Partially reconstructed information has limited value. This is due to the layered nature of Stratis. It's not sufficient to know some subset of the device mapping levels. Since they are layered, recovering e.g. L0-L2 layouts allows no data to be recovered without also knowing how L3 and L4 are mapped on top, and vice versa.
3. L0-L4 metadata should require relatively few updates per day, since the changes it would reflect are blockdevs being added/removed from the pool, or thinpool data device expansions. Infrequent updates reduces the likelihood of corruption¹⁹.

L5-L7 is stored on the Metadata Volume on an XFS filesystem. Partial data recovery of that information is possible.

In addition to Stratis-specific metadata, device-mapper layers such as cache, raid, thin, as well as XFS filesystems, all have their own metadata. Stratis would rely on running each of their specific repair/fsck tools in case they reported errors.

9.4 Cache Tier

The Cache Tier is a secondary optional stack that, if present, sits underneath the Flex Layer, and above the

¹⁹citation needed?

9.4.1 Requirements

1. Caching may be configured for redundancy, or no redundancy.
2. Caching may be configured for write-back and write-through modes.
3. Stratis concatenates all cache blockdevs and uses the resulting device to cache the thinpool device (L4). This lets all filesystems benefit from the cache.
4. Cache blocksize should match thinpool datablock size.
5. Removing cache tier comes with performance hit and “rewarming” phase
6. For write-back caching, Cache tier must be redundant if data tier is redundant.

9.5 Cache Tier Metadata

9.5.1 Cache Tier Metadata Requirements

1. Identify all blockdevs that are part of the pool’s cache tier, the configured redundancy level, and other cache-specific configuration parameters (e.g. WT/WB, block size, cache policy)
2. Cache tier supports up to 8 devices.

10 Implementation Details

10.1 'stratis' command-line tool

Stratis' command-line tool is currently written in Python. Since this is only used after the system is booted by the administrator, Python's interpreted nature and overhead is not a concern.

10.2 stratisd

Stratisd needs to be implemented in a compiled language, in order to meet the requirement that it operate in a preboot environment. A small runtime memory footprint is also important.

stratisd is written in Rust. The key features of Rust that make it a good choice for stratisd are:

- Compiled with minimal runtime (no GC)
- Memory safety, speed, and concurrency
- Strong stdlib, including collections
- Error handling
- Libraries available for Dbus, device-mapper, JSON serialization, and CRC
- FFI to C libs if needed
- Will be available on RHEL 7 in delivery timeframe; currently packaged in Fedora

Other alternatives considered were C and C++. Rust was preferred over them for increased memory safety and productivity reasons.

10.3 device-mapper names

If `stratisd` terminates unexpectedly and is restarted, it needs to rebuild its knowledge of the running system. This includes not only re-enumerating blockdevs to find Stratis pool members, but also determining the current state of the device-mapper targets that make up pools. A restarting `stratisd` needs to handle if none, some, or all of the expected DM devices are present, and if present DM devices are working correctly, or in an error state.

To these ends, Stratis uses consistent naming for device-mapper targets. This lets `stratisd` more easily determine if DM devices already exist, and avoids leaking old DM mappings.

10.3.1 Naming convention Requirements

- Globally unique
- Maximum 127 characters
- Differentiate between Stratis and other DM devices
- Forward-compatible to allow `Stratisd` updates
- Human-readable
- Easily parsable

10.3.2 Naming Convention

Stratis DM names consist of five required and two optional parts, separated by a '-'.

Part	Name	Max length	Required	Description
1	stratis-id	7	y	Universal DM differentiator: 'stratis'
2	format-version	1	y	Naming convention version: '1'
3	private	7	n	Optional indicator 'private'
4	pool-id	32	y	ASCII hex UUID of the associated pool
5	layer-name	16	y	Name of the Stratis layer this device is in
6	layer-role	16	y	Name of the role of the device within the layer
7	role-unique-id	40	n	Role-specific unique differentiator between multiple devices within the layer with the same role

- The maximum length (adding 6 '-'s as separator) is 125, to stay within the DM name limit of 127 characters.
- "private" is included in names for DM devices that are internal and that should be excluded from content scanning by tools such as `blkid`.
- Characters for each part are drawn solely from the character classes '[a-z]' and '[0-9]' except that part 7 may also use the '-' character. These restrictions meet D-Bus and `udev` requirements²⁰.)

10.3.3 Per-Layer Naming Definitions

Layer 0: Blockdev

Not needed.

Layer 1: Integrity

TBD

²⁰See `libdm/libdm-common.c _is_whitelisted_char()` in the `lvm2` code for more.

Layer 2: Redundancy

TBD

Layer 3: Flex

Private	layer-name	layer-role	role-unique-id	Description
y	flex	thinmeta	(none)	Linear device for thinpool metadata
y	flex	thindata	(none)	Linear device for thinpool data
y	flex	mdv	(none)	Linear device for Metadata Volume (MDV)

Example: stratis-1-private-2836fca3047bba22938475638abcd840-flex-thinmeta

Layer 4: Thinpool

Private	layer-name	layer-role	role-unique-id	Description
y	thinpool	pool	(none)	Thinpool device

Example: stratis-1-private-2836fca3047bba22938475638abcd840-thinpool-pool

Layer 5: Thin volume

Private	layer-name	layer-role	role-unique-id	Description
n	thin	fs	32-char ASCII hex fs UUID	Thin volume

Example: stratis-1-2836fca3047bba22938475638abcd840-thin-fs-ca7efca3047bba22938475638abc3141

Layer 6: Encryption

TBD

10.4 device-mapper minimum version

Stratisd device-mapper minor version 37 or greater, for DM event poll() support and support for event_nr in list_devices ioctl.

10.5 OS Integration: Boot and initrd

Since we want to allow Stratis to be used for system files, Stratis needs to run in the initrd preboot environment. This allows it to activate pools and filesystems so that they can be mounted and accessible during the transition to the main phase of operation.

The use of D-Bus is not possible in the preboot environment. Therefore, Stratis must tolerate its absence.

The most likely implementation of these requirements is a “boot-init” mode for Stratis, in which it discovers blockdevs and activates pools, but does not connect to D-Bus, and then exits. Then, as part of the systems main phase, Stratis will run again and stay running, in order to handle pool monitoring and commands over D-Bus, as normal.

10.6 OS Integration: udev

The udev library “libudev” enables access to the udev device database. This allows library users to enumerate block devices on the system, and includes attributes describing their contents, such as what filesystem or volume manager signature was detected. (libudev uses libblkid for this, which recently had Stratis signature support added.) The primary benefit of this is to perform the time-consuming block device scan only once, and to alleviate library users from interpreting block device contents.

On boot, Stratis uses libudev to enumerate Stratis block devices on the system, reads the Stratis metadata from each, and activates pools that are complete. Later, during the main running phase, Stratis monitors udev events for newly-added block devices, so that if missing Stratis pool members are connected to complete a pool, the pool can be activated and used.

10.7 OS Integration: /dev entries

Stratis allows the user to create filesystems, which then can be mounted and used via `mount(8)` and the `fstab(5)`. `Stratisd` creates a `/stratis` directory. It creates `/stratis/<poolname>` for each pool present on the system, and `/stratis/<poolname>/<filesystemname>` for each filesystem within the pool. Changes such as creations, removals, and renames are reflected in the entries under `/stratis`. These entries give the user a well-known path to a device to use for mounting the Stratis filesystem. Filesystems may also be listed in `/etc/fstab` using XFS UUID.

10.8 Partial Pool Activation

Stratis must handle if some blockdevs that make up a pool are not present. Its initial policy is not to activate a pool until all members are present. However, there are scenarios where this may not be strictly necessary, such as a blockdev that has not had any Flex Layer allocations on it, or (in Stratis 2.0) if redundancy allows the RAID Layer to function in a degraded state. The initial policy will likely be modified over the course of development to accommodate these cases.

10.9 Snapshots

Stratis's current snapshot implementation is characterized by a few traits:

- A snapshot and its origin are not linked in lifetime. i.e. a snapshotted filesystem may live longer than the filesystem it was created from.
- A snapshot of a filesystem is another filesystem.
- A filesystem may be snapshotted while it is mounted or unmounted.
- Each snapshot uses around half a gigabyte of actual backing storage, which is needed for the XFS filesystem's log.

These may change in the future.

10.10 Backstore Internals

The backstore is divided into two tiers: the data tier, and an optional cache tier. Each tier has its own set of physical block devices. The goal of each tier is to provide a single linear device that the flex layer (or another tier) can easily build on top of.

A tier is created with a certain feature set, which results in an internal layering of devices as needed to support those features. Features such as redundancy, integrity, and dedupe/compression. The features a tier supports are fixed at tier creation time. However, the block devices that make up the tier may change. A blockdev may increase in capacity, or new blockdevs may be added.

This requires each tier to support:

- Today
 - `add_blockdev` (add a new blockdev to the tier)
 - `blockdevs` (list/iterate)
- Near Future

- extend_blockdev (start using added space on a blockdev that has become larger)
- Far future
 - prepare_to_remove_blockdev (move data off a blockdev in preparation for its removal, if space is available elsewhere)
 - remove_blockdev (remove a blockdev that has no data on it)

The tier includes optional internal support for multiple features, which also are implemented using DM devices.

At the “bottom” of the tier are blockdevs. These blockdevs are mapped through layers that add value, such as integrity or compression.

Each layer takes a list of blockdevs and converts it to a list of “better” blockdevs, whose total size is likely different. (Integrity and raid will be smaller, vdo will present a single much larger block device, etc.)

While each intermediate layer may provide an array of blockdevs, the “cap” layer of the tier presents a single linear blockdev that maintains the location of each presented block and never shrinks, and hides the interior complexity of the tier from upper users.

The ordering of layers (from bottom to top) within a tier is:

1. blockdev

Blockdevs supply available space to the tier. Blockdevs may grow, for example if they happen to be further virtualized on top of a storage appliance, Amazon EBS, or Ceph block storage.
2. integrity

When integrity=true, each blockdev’s usable space is mapped to the dm-integrity target, which will slightly reduce usable space, and cause a performance hit, in order to detect corruption on reads. It is believed that integrity devices cannot change size. Therefore, extending a blockdev may require a second dm-integrity device be created.
3. raid

when raid=true, Stratis uses the dm-raid target to create software RAID in top of the devices presented to it (These may be either “raw” blockdevs or integrity-enabled blockdevs.) In order to maximize flexibility, Stratis creates multiple raidsets by defining a maximum for each raid member’s size, as well as the number of members in a raidset. It also may choose a raid member size smaller than the maximum, in order to ensure between 4-8 raidsets are created.

if additional block resources are made available, the raid layer constructs additional raidsets.
4. vdo

When vdo=true, the vdo layer will concatenate the blockdevs it is given (if needed) of any previously-mentioned type and create a single VDO device. The layer must manage extending both the logical size (the top) in response to space requests from above, as well as requesting additional space from lower layers to extend the physical size (the bottom) if needed.
5. cap

If presented with more than one blockdev, or the blockdev has a nonzero offset, the “cap” layer will ensure the Tier presents a single blockdev with consistent block mapping for use above the tier by creating a Linear device that never relocates previously-mapped block ranges.

10.10.1 Demand-based allocations

Layers should not consume the entire space available to them when constructing devices, but instead grow existing mapped allocations (or create new ones) as the total demands of upper layers grow larger. This is preferred over a “greedy” strategy because any mapping that is created and has data on it must be maintained for the life of the pool. There’s no good way to tell when blocks in an existing device mapping become used. Starting small and then growing allocations gives each layer a way to track regions that can be reclaimed if needed. This will make eventual support for blockdev removal easier.

Another reason is that space requirements may expand or shrink as a request passes down the layers. For example, VDO may satisfy a request for more logical space without any additional lower-level space needed. But, a small request to the RAID layer may require the construction of an entirely new raidset with a large minimum size from unused space, if the layer’s existing capacity is full.

10.10.2 Common Layer API

TBD

10.11 Operation States

When encountering errors, Stratis must handle them if possible, but there are also errors that are severe enough to hamper Stratis’s ability to function. When these occur, instead of terminating, Stratis continues by transitioning to a less-capable operation state. This allows some measure of continued monitoring and enables its condition to be visible to the user through the API.

Stratis has operation states for the pools, meta and data extend, and also blockdevs (listed on page ??).

10.11.1 Pool States

In addition to the main states a pool may be in, a Running pool may also be Complete or Incomplete, and a Stopped pool may have one of several things preventing it from running.

State	Description	Transitions To
Initializing	Pool startup In progress	Running, OutOfDataSpace, ReadOnly, Failed
Running	Pool is read/write - normal operation state	OutOfDataSpace, ReadOnly, Failed, Stopping
OutOfDataSpace	Meta data may be changed - new data may not be allocated. Data device is update only	Running, ReadOnly, Failed, Stopping
ReadOnly	Meta data may not be changed - data is read only	Running, OutOfDataSpace, Failed, Stopping
Failed	All IO is failed	Running, OutOfDataSpace, ReadOnly, Stopping
Stopping	Pool is in the teardown process	

10.11.2 Pool Extend States

As filesystems write data, the meta and data devices backing the pool may need to be extended to accomidate growing pool.

State	Description	Transitions To
Initializing	Pool startup In progress	Good, DataFailed, MetaFailed, MetaAndDataFailed
Good	Normal operation	DataFailed, MetaFailed, MetaAndDataFailed
DataFailed	Extend of the pool data device failed	Good, MetaAndDataFailed
MetaFailed	Extend of the pool meta device failed	Good, MetaAndDataFailed
MetaAndDataFailed	Extend of both the meta and data devices failed	Good, DataFailed, MetaFailed

10.11.3 Pool Space States

Pool space available also has several states it can be in. The pool may transition freely between these. Being in a Critical state can result in the Pool State transitioning to Stopped/OutOfSpace.

State	Description	Action
Good	Space available is above Low threshold	None
Low	Space available is below Low threshold but above Critical	Alert the user, run maintenance operations
Critical	Space available below Critical threshold	Alert the user of critical state, run maintenance operations

10.12 Licensing

Currently, stratisd and all Rust libraries included in the stratis-storage organization meant for stratisd are licensed as MPLv2. This conclusion was reached after careful consideration of compatibility with other licenses in the open source ecosystem. MPLv2 allows some major benefits:

1. It is compatible with more permissive licenses in the open source ecosystem such as BSD and MIT licenses. Given the prevalence of more permissive licenses in the Rust ecosystem, this is an important consideration.
2. It is compatible with GPL code. This permits stratisd to incorporate GPL code and become effectively GPL without ever changing the MPLv2 license.
3. If GPL code is incompatible with the license of a dependency added later, MPLv2 allows the removal of GPL code and migration of the GPL code into a separate service to permit the addition of the conflicting dependency. This can all be done without relicensing.

After evaluating the options, the MPLv2 seemed to be the most flexible license that still fulfilled the requirements for this project.

Part IV

Development Plan

11 Delivery of Features

11.1 Stratis version 0.1

Simplest thing that does something useful

1. Create a pool

2. Destroy a pool
3. Create a filesystem
4. Destroy a filesystem
5. List filesystems
6. Rename filesystems
7. List pools
8. Rename pools
9. List blockdevs in a pool
10. Redundancy level: none
11. D-Bus API
12. Command-line tool
13. Save/restore configuration across reboot
14. Initial disk labeling and on-disk metadata format
15. Error handling for missing, corrupted, or duplicate blockdevs in a pool
16. thin/cache metadata validation/check (i.e. call `thin_check` & `cache_check`)

11.2 Stratis version 0.5

Add cache tier & basic snapshot support

1. List cache blockdevs in a pool
2. Add cache blockdevs
3. Remove cache blockdevs
4. Write-through caching only
5. Create/destroy primary snapshots
6. Working D-Bus signals

11.3 Stratis version 1.0

Minimum Viable Product

1. Snapshot management: auto snaps, date-based culling, “promotion” from snap to “primary”
2. Monitor pool(s) for getting close to capacity, and do something (remount ro?) if dangerously full
3. Notification method to the user if pool is approaching user or system-defined capacity
4. Maintain filesystems: Grow a filesystem as it nears capacity
5. Maintain filesystems: Run `fstrim` periodically to release unused areas back to thinpool
6. Add and use an additional blockdev

11.4 Stratis version 2.0

Add Redundancy Support

1. Remove an existing blockdev
2. Redundancy level: raid1
3. Redundancy level: raid5
4. Redundancy level: raid6
5. Redundancy level: raid10
6. Cache redundancy level: raid1
7. Write-through caching enabled
8. Quotas
9. Blockdev resize (larger)
10. Spares

11.5 Stratis version 3.0

Rough ZFS feature parity. New DM features needed.

1. Send/Receive
2. Integrity checking (w/ self-healing only if on raid)
3. Raid scrub
4. Compression
5. Encryption
6. Dedupe
7. Raid write log (on ssd? To eliminate raid write hole)

11.6 Stratis version 4.0

Future features and evolution

1. Change a pool's redundancy level
2. Boot from a filesystem
3. Libstoragemgmt integration
4. Multipath integration
5. Tag-based blockdev and filesystem classification/grouping
6. Mirroring across partitions within a pool, for multi-site or across hw failure domains (shelves/racks)
7. Support for byte-addressible persistent memory

References

- [1] Anne Mulhern. justbytes documentation. <http://pythonhosted.org/justbytes/>.
- [2] Anne Mulhern. The computation and representation of address ranges: With an introduction to the python justbytes library. https://mulhern.fedorapeople.org/justbytes_presentation.pdf, August 2016.

Appendices

A Initial options for adapting existing solutions

As part of early requirements-gathering, the team looked at existing projects in this space, both as candidates for building upon to create a solution, as well as if an existing project could be extended to meet the requirements.

A.1 Extending an existing project

A.1.1 SSM

System Storage Manager (SSM) provides a command line interface to manage storage in existing technologies. Our interest in SSM was to determine if it would be an existing project we could extend to meet our requirements.

SSM provides a unified interface for three different “backends”: LVM, Btrfs, and crypto. However, if we wish to provide a simple, unified experience, the first step would likely be to pick one of the backends and build around its capabilities. This eliminates complexity from the CLI -- no need for the user to pick a backend or encounter commands that happen to not work based upon the chosen backend, but obviates much of the point of SSM.

SSM does not provide a programmatic API. It internally contains “ssmlib”, which could be enhanced and exposed, but would be Python-only. ssmlib is also built around executing command-line tools, which can cause issues.

SSM is not a daemon. We’d need to modify SSM to operate on a daemon model. An ongoing presence is needed for fault monitoring but also automatic filesystem and thinpool extensions.

SSM doesn’t currently support RAID5/6, thin provisioning, or configuring a cache tier.

SSM is written in Python, which would limit its ability to be used in an early-boot environment.

SSM does not provide functionality for error recovery. If the storage stack encounters an error the user has to use the individual tools in the stack to correct. Thus greatly diminishing the ease of use aspect and value proposition of SSM.

Analysis

Extending SSM does not meet the requirements.

A.1.2 LVM2

Logical Volume Manager (LVM2) is the nearly universally-used volume manager on Linux. It provides the “policy” that controls device-mapper. It adds:

- On-disk metadata format to save and restore configuration across boot
- Usage model built on Physical Volume, Volume Group, and Logical Volume (PV, VG, LV) concepts.

- A comprehensive set of command line tools for configuring linear, raid, thinpool, cache, and other device-mapper capabilities
- Monitoring, error handling, and recovery
- LV resize; PVs may be added or removed from a VG
- Snapshots and thin snapshots
- Choice of user-guided or automatic allocation/layout within the VG

Analysis

Adding the capability to manage filesystems to LVM isn't something that has been much considered. Extending LVM2 would make it very hard to achieve simplicity of interface, given the conflicting requirement to maintain backwards compatibility with what LVM provides now.

A.2 Building upon existing projects

A.2.1 XFS

XFS is a highly respected non-volume-managing filesystem. To meet the goal of eliminating manual filesystem resizing by the user, Stratis requires the filesystem used have online resize (or at least online grow) capabilities, which XFS does. In the absence of online shrink, Stratis would rely on trim to reclaim space from an enlarged but mostly empty filesystem, and return it to the thin pool for use by other filesystems.

Use of XFS on top of thin provisioning also makes proper initial sizing important, as well as choosing sizes for XFS on-disk allocations that match those used by the underlying thin-provisioning layer, to ensure behavior with the two layers is optimal.

Analysis

XFS meets the requirements and currently seems like the best choice.

A.2.2 device-mapper

device-mapper is a framework provided by the Linux kernel for creating enhanced-functionality block devices on top of physical block devices. These new devices can add support for RAID, thin provisioning, encryption, multipath devices, caching devices, and more. The framework provides the ability to configure and layer these capabilities, but no facilities for saving or restoring a configuration. device-mapper provides mechanism, but no policy.

Analysis

Using devicemapper directly would require that an upper layer implement its own on-disk metadata format and handle some tasks in a similar manner to LVM2.

A.2.3 LVM2

See A.1.2 for a description of LVM2 capabilities.

LVM is a mature software project that implements volume management. For Stratis, the question is whether the benefits of internally using LVM2 outweigh the costs.

Issues with Building on LVM2

Note: This assumes the implementation described in Part III.

Note: lvm-team has raised objections to items on this list.

- Policy+mechanism vs policy+policy+mechanism: LVM2 is configurable but has limitations. e.g. we might wish to let the user define a block device as only to be used to replace a failed disk in a raidset. However LVM `raid_fault_policy="allocate"` will use *any* free PV, not just one explicitly reserved.
- A good API needs the ability to convey meaningful and consistent errors for other applications to interpret. The lvm command line employs a simple exit code strategy. The error reason is embedded in stderr in free form text that changes without notice. Thus it is virtually impossible for any lvm command line wrapper to provide meaningful and consistent error codes other than success or failure. Note: Lvm has recently added JSON output which contains the ability to add more meaningful and useful error codes, but this functionality is not implemented and non-trivial in scope to complete.
- lvm-dbus cannot be used because it requires Python and D-Bus, neither of which are available in initrd
- Stratis-managed LVM2 devices would show up in LVM2 device & volume listings, which would cause user confusion
- Using LVM2 for metadata tracking is good, but only if upper layer has no metadata storage needs of its own. What about tags? Tags can't store JSON objects since '`[]{},"`' are not allowed in tags.
- LVM2 metadata format prevents new metadata schemes, such as tracking thin volumes separately from PV metadata, or metadata backup copy not also at the tail of the blockdev.
- Use of new device-mapper features delayed by LVM2 implementation and release cycle.
- One big argument by LVM2 proponents is that LVM2 is a large, long-lived project that has learned many things the hard way, and it would be foolish to abandon all that value by starting over.
 - Must we use the code, or can we take lessons from LVM2 devs and incorporate them independently? Maybe fix some things that backwards-compatibility makes impossible to fix in LVM2?
 - Large parts of the codebase don't benefit Stratis:
 - * File-based & configurable Locking: not needed since everything is serialized through `stratisd`
 - * daemons/* including `clvmd`
 - * Udev: `stratisd` assumes udev & listens for udev events
 - * Filter/global_filter
 - * Caching: not needed, daemon is authoritative
 - * profiles
 - * preferred names ordering
 - * `lvm.conf` display settings: not needed, up to API client
 - * dev_manager: Stratis layers are predefined, much simpler
 - * config_tree
 - * report: beyond Stratis scope
 - * Command-line tools, option parsing: handled in cli, reduced in scope
 - * lib/misc/*: not needed or handled via libraries

- * Multi-metadata-format support
- What would Stratis benefit from?
 - * on-disk metadata format
 - * Best policy for duplicate/absent/corrupted block devices
 - * fault tolerance/recovery
 - * pool/snapshot monitoring

Analysis

While we cannot dismiss using LVM as an option for the future, currently there are some areas that it does not meet Stratis requirements. There are also questions about the best way to interface with LVM that need to be resolved prior to its adoption.

A.3 Conclusions

Based on looking at the existing available building blocks, the best option is to build Stratis as a new project that initially makes use of XFS and device-mapper in its implementation. In parallel, request enhancements to LVM2 to enable its substitution for device-mapper when the enhancements are implemented. This lets Stratis proceed without delay to a point where it can be placed in prospective users' hands to start getting feedback, and will allow Stratis to eventually use LVM2, and avoid duplicating functionality that LVM2 already provides.